

Hardware-Software Co-Design of a 4-bit Quantized MLP

From PyTorch Training to FPGA RTL

Deployment

Presenter: Youssef Mohamed

Project Overview & Objectives

Goal

Design an FPGA-based accelerator for Handwritten Digit Recognition (MNIST).


Key Achievement

Successfully verified a bit-exact match between the Python Quantized model and Verilog RTL simulation.

Project Scope

- **Training:** Develop a hardware-friendly Neural Network in PyTorch.
- **Quantization:** Compress model from Float32 to 4-bit Integers (Int4).
- **Hardware:** Implement the accelerator in Verilog RTL (simulation only).

Model Architecture (Software Design)



Network Structure: 3-Layer MLP

- **Input:** 784 neurons (28×28)
- L1: 64 Neurons + ReLU
- L2: 32 Neurons + ReLU
- **Output:** 10 Neurons (Digits 0-9)

Design Decision: Bias = False

Hardware forced: $y = Wx$

Hardware Benefit: Eliminates extra Adders and Muxing logic.

Impact: Negligible accuracy drop (<0.5%).

Quantization Methodology

⌚ Method

Post-Training Static Symmetric Quantization.

Bit-Width: 4-bit Weights and Activations.

⚖️ Why Symmetric?

Maps Floating-point to Integer .

Equation: $Y = W * X$

Avoids expensive zero-point logic:

$$Y = (X - ZX)(W - ZW)$$

📊 Data Ranges

Weights: Signed 4-bit
(-8 to +7)

Activations: Unsigned 4-bit
(0 to 15) via ReLU

Design Philosophy: Serial Reuse


Instead of massive parallelism (area-expensive), we use a single Calculation Core reused sequentially.

Memory Hierarchy

- **ROM:** Stores Weights (linear addressing).
- **RAM:** Stores Input Image and Intermediate Layer Outputs.

Control Logic

- **Muxing:** Data Router switches inputs based on layer state.
- **FSM:** Orchestrates data movement.



Finite State Machine (FSM)

1. LAYER_X

Fetch data → Accumulate
(Repeat 784 times)

2. WRITE

Store result to RAM
(activates wen signal)

3. INC

Reset Accumulator
Increment Neuron Counter

4. ARGMAX

Layer 3 Only
Tracks max_logit for prediction

Challenge 1: Signed/Unsigned Arithmetic

⚠ The Issue

- Inputs are **Unsigned** (0..15).
- Weights are **Signed** (-8..7).
- Verilog standard multiplication can misinterpret the Unsigned input as a negative 2's complement number if the MSB is 1 (e.g., 1111 = 15 or -1?).



```
assign product = (current_input === 4'bx || current_weight === 4'bx)
? 0 : $signed({1'b0, current_input}) * $signed(current_weight);
```

✓ The Solution

Explicitly zero-padded the input to force positive interpretation.

Mathematically correct mixed-sign multiplication.

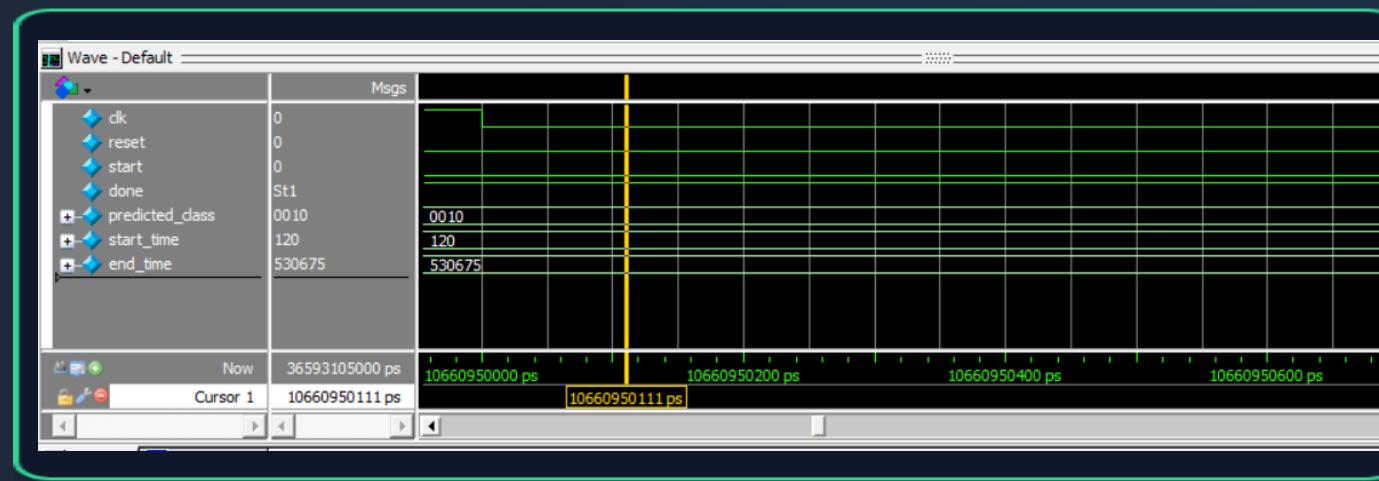
Challenge 2: Memory Latency & Pipelining

The Issue: BRAM Latency

Block RAM has a 1-cycle read latency. Accumulating data_out in the same cycle as addr request reads "garbage."

The Solution: 2-Stage Pipeline

- ❖ Cycle 0: pipeline_valid <= 1 (Request)
- ❖ Cycle 1: pipeline_valid_d <= 1 (Accumulate)



Result: Perfect synchronization; the accumulator only updates when valid data arrives.

Challenge 3: The "Epoch Trap"

Observation

Extended training (30 epochs) caused Quantized Accuracy to crash to **73.32%**.

Root Cause: Outlier Weights

As training continues, the model pushes some weights to extreme values (e.g., -3.5 or +4.0). Symmetric Quantization scales the entire 4-bit range based on the largest number, wiping out precision for small weights.

Solution: Used 6 EPOCHS ONLY

```
Epoch 25/30 complete. Loss: 0.0016
Epoch 26/30 complete. Loss: 0.0001
Epoch 27/30 complete. Loss: 0.0016
Epoch 28/30 complete. Loss: 0.0062
Epoch 29/30 complete. Loss: 0.0001
Epoch 30/30 complete. Loss: 0.0008
```

```
--- Calculating Original (Floating-Point) Model Accuracy ---
Original Model Accuracy: 97.41%
```

```
--- Starting Quantization ---
Saved w1.mem
Saved w2.mem
Saved w3.mem
Saved input1.mem
```

```
Expected Label for input.mem: 1
```

```
--- Golden Reference (Integer Simulation) ---
Layer 3 Output (logits): [-10.  8. -5. -6.  4. -9. -10.  6. -5. -1.]
Predicted Digit: 1
```

```
--- Calculating Quantized Model Accuracy (Hardware Simulation) ---
Quantized Model Accuracy: 73.32%
```

Final Results & Verification

```
... Training for 6 epochs (bias=False)...  
Epoch 1/6 complete. Loss: 0.3354  
Epoch 2/6 complete. Loss: 0.0656  
Epoch 3/6 complete. Loss: 0.3705  
Epoch 4/6 complete. Loss: 0.0265  
Epoch 5/6 complete. Loss: 0.0149  
Epoch 6/6 complete. Loss: 0.0371  
  
--- Calculating Original (Floating-Point) Model Accuracy ---  
Original Model Accuracy: 97.09%  
  
--- Starting Quantization ---  
Saved w1.mem  
Saved w2.mem  
Saved w3.mem  
Saved input1.mem  
  
Expected Label for input.mem: 1  
  
--- Golden Reference (Integer Simulation) ---  
Layer 3 Output (logits): [-17. 32. -17. -16. -13. -23. -15. 8. -10. -5.]  
Predicted Digit: 1  
  
--- Calculating Quantized Model Accuracy (Hardware Simulation) ---  
Quantized Model Accuracy: 95.56%
```

```
# DEBUG: L3 Neuron 0 | Logit: 1 | Current Max: -262144  
# DEBUG: L3 Neuron 1 | Logit: -34 | Current Max: 1  
# DEBUG: L3 Neuron 2 | Logit: 59 | Current Max: 1  
# DEBUG: L3 Neuron 3 | Logit: 11 | Current Max: 59  
# DEBUG: L3 Neuron 4 | Logit: -89 | Current Max: 59  
# DEBUG: L3 Neuron 5 | Logit: -42 | Current Max: 59  
# DEBUG: L3 Neuron 6 | Logit: -57 | Current Max: 59  
# DEBUG: L3 Neuron 7 | Logit: 6 | Current Max: 59  
# DEBUG: L3 Neuron 8 | Logit: 17 | Current Max: 59  
# DEBUG: L3 Neuron 9 | Logit: -37 | Current Max: 59  
# -----  
# Inference Complete.  
# Predicted Class: 2  
# Total Cycles: 53055  
# Total Time: 530555 ns
```


... Saved input1.mem
Saved input1.mem for Image Index 444 (True Label: 2)

--- Python Golden Prediction ---
Logits: [-10. 4. 53. 20. -59. -14. -50. 18. 18. -7.]
Predicted Class: 2

Final Quantized Model (6 Epochs): 95.56% Accuracy

Hardware output exactly matches the Python Golden Reference.

Simulation Results & Verification



✓ **Verification Successful**

RTL Simulation achieved a **Bit-Exact Match** to the Quantized Python model.

Testbench: tb_mlp.v with w1.mem and input1.mem

Q & A

Thank you for your attention.